

CS415 Module 2 Part C - Process Control

Athens State University

Contents

1	The <code>wait()</code> System Call	1
2	Waiting for a specific process	2
3	Orphans and Zombies	4
4	Key Points	5

1 The `wait()` System Call

It's a long wait

- We noted in our discussion about `fork()` and `exec()` that most cases require that the parent keep track of its children to find out when and how they terminate

```
1 #include <sys/wait.h>
   pid_t wait(int *status)
```

What's happening with `wait()`

- If no previously waited-upon child of the calling process has yet to terminate, block the caller until one of the children terminate.
- If any child has terminated at the time of the call, immediately return to the calling process
- If the `status` pointer is not NULL, information about what caused the child to terminate is placed into the integer pointed to by `status`.
- The kernel update process CPU and resource usage stats to running totals for all of the children of this parent process
- The return type of the system call is the process ID of the child that terminated.
- Upon error, the system call returns `-1` to the caller

Note that this means that we use the following loop as a means to force the parent to wait until all of its children have terminated:

```

pid_t childPid;
child_pid = wait(NULL);
while (child_pid == -1) {
    child_pid = wait(NULL);
}
if (errno != ECHILD) {
    errExit("wait");
}

```

Using wait()

```

for (j = 1; j < argc; j++)
{
    /* Create one child for each argument */
    switch (fork()) {
        case -1:
            errExit("fork");
        case 0: /* Child sleeps for a while then exits */
            printf("[%s] child %d started with PID %d, sleeping %s "
                    "seconds\n", currTime("%T"), j, (long) getpid(), argv[j]);
            sleep(getInt(argv[j], GN_NONNEG, "sleep-time"));
            _exit(EXIT_SUCCESS);
        default: /* Parent just continues around loop */
            break;
    }
}

```

Using wait()

```

1 numDead = 0;
for (;;) { /* Parent waits for each child to exit */
    childPid = wait(NULL);
    if (childPid == -1) {
        if (errno == ECHILD) {
            printf("No more children - bye!\n");
            exit(EXIT_SUCCESS);
        } else { /* Some other (unexpected) error */
            errExit("wait");
        }
    }
}

13 numDead++;
printf("[%s] wait() returned child PID %d (numDead=%d)\n",
15 currTime("%T"), (long) childPid, numDead);
}

```

2 Waiting for a specific process

Limitations of the wait() system call

- If a process creates more than one child, the wait() does not allow us to wait for termination of a specific child

- If no child has yet to terminate, then `wait()` always blocks.
 - There are use cases where it's better to just say that no children have terminated
 - We don't always want to block; i.e., stop everything
- The `wait()` system call can tell you what's happening with children that have terminated

The `waitforpid()` system call

```
#include <sys/wait.h>
pid_t waitpid(pid_t pid, int *status, int options)
```

- Returns the process ID of the child, 0, or -1 (on error)

The meaning of `pid` variable in `waitforpid()`

- If `pid` is greater than 0, wait for the child whose process ID is equal to `pid`
- If `pid` is equal to 0, wait for any child in the same *process group* as the parent.
- If `pid` is less than -1, wait for any child whose process group id equals the absolute value of `pid`
- If `pid` is equal to -1, wait for any child

Note that the call `wait(&status)` is equivalent to calling `waitpid(-1, &status, 0)`.

The meaning of the `options` variable in `waitforpid()`

- The `options` argument is a bitmask that you can OR with one or more of the following flags:
 - `WUNTRACED`: Include status information about whether a process terminated or stopped
 - `WCONTINUED`: Include status information about children who have been restarted by a signal
 - `WNOHANG`: Do not block, always return 0

The `options` variable is important to us due to the `WNOHANG` bitmask. Setting this option converts the system call from blocking to polling, which means that we will not wait until child terminates or stops.

Interpreting the status from `wait()` and `waitforpid()`

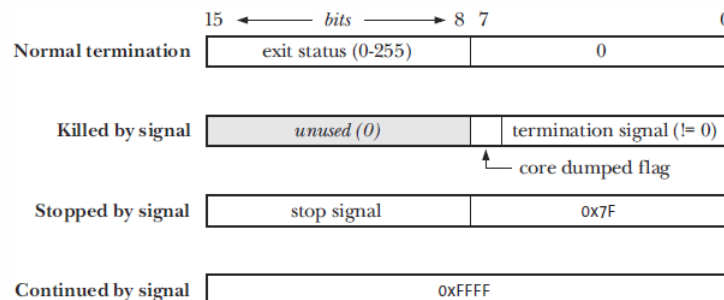


Figure 26-1: Value returned in the `status` argument of `wait()` and `waitpid()`

Fortunately, the header file for the `wait()` and `waitforpid()` system calls provide C macros that parse status codes:

`WIFEXITED(status)` Return true if the child process exited normally

`WEXITSTATUS(status)` Returns the exit status of the child process

`WIFSIGNALED(status)` Returns true if the child process was called by a signal

`WIFSTOPPED(status)` Returns true if the child process was stopped by a signal

```
1 #include <sys/wait.h>
2 #include "print_wait_status.h" /* Declares printWaitStatus() */
3 #include "tldpi_hdr.h"
4 int main(int argc, char *argv[])
5 {
6     int status;
7     pid_t childPid;
8     if (argc > 1 && strcmp(argv[1], "--help") == 0)
9         usageErr("%s [exit-status]\n", argv[0]);
10    switch (fork()) {
11        case -1: errExit("fork");
12        case 0: /* Child: either exits immediately with given
13                status or loops waiting for signals */
14            printf("Child started with PID = %ld\n", (long) getpid());
15            if (argc > 1) /* Status supplied on command line? */
16                exit(getInt(argv[1], 0, "exit-status"));
17            else
18                for (;;) /* Otherwise, wait for signals */
19                    pause();
20            exit(EXIT_FAILURE); /* Not reached, but good
                                practice */
    }
```

```
1     default: /* Parent: repeatedly wait on child until it
2              either exits or is terminated by a signal */
3         for (;;) {
4             childPid = waitpid(-1, &status, WUNTRACED);
5             #ifdef WCONTINUED /* Not present on older versions of Linux */
6                 | WCONTINUED
7             #endif
8             );
9             if (childPid == -1)
10                errExit("waitpid");
11            /* Print status in hex, and as separate decimal bytes */
12            printf("waitpid() returned: PID=%ld; status=0x%04x (%d,%d)\n",
13                (long) childPid,
14                (unsigned int) status, status >> 8, status & 0xff);
15            printWaitStatus(NULL, status);
16            if (WIFEXITED(status) || WIFSIGNALED(status))
17                exit(EXIT_SUCCESS);
18        }
19    }
```

3 Orphans and Zombies

Orphaned Children

- A process is *orphaned* when its parent terminates before its children
- In this case, the orphaned child is adopted by the system `init` process
- The orphan's parent process id is set to 1. This provides a means for one to determine if a child's parent is still alive

Zombie Processes

- What happens to a child that terminates before its parent has a chance to perform a `wait()`?
- The child may have finished, but the parent should still be permitted to perform a `wait()` at some later time to determine what happened
- The kernel addresses this question by turning the child into a *zombie process*
 - The kernel keeps an entry for terminated process in its process table
 - The downside is that if the parent never calls `wait()` then the entry is maintained indefinitely in the process table
 - Killing zombie processes require you to kill their parent (like an episode of *Walking Dead*, Linux style)

4 Key Points

Key Points

- The `wait()` and `waitforpid()` system calls
- Orphans and zombie processes