

CS415 Module 5 Part B - IPC, Shared Memory, and Named Pipes

Athens State University

Outline

Contents

1	Shared Memory	1
1.1	Case Study: Data Transfer Between Processes	2
1.2	POSIX Shared Memory	3
2	Pipes, Named and Otherwise	8
2.1	Named Pipes	10

1 Shared Memory

System V vs. POSIX Shared Memory

- There are two different shared memory interfaces provided in Linux
 - The older UNIX System V interface
 - The more recent POSIX standard interface
- There are reasons why one might use one or the other interface
 - We will start with the System V interface


The Six Steps Required to Use SysV Shared Memory

- Call the `shmget()` system call create a new segment or connect to an existing segment
- Use `shmat()` to *attach* the shared memory segment to a memory segment in the calling process's address space
- Use the memory location as if it was in the local address space
- Call the `shmdt()` system call to detach the shared memory from the process's address space
- Call `shmctl()` to delete the shared memory segment. The OS will not actually destroy the segment until all users are done with it

The Critical System Calls for SysV Shared Memory

```
1 #include <sys/types.h>
2 #include <sys/shm.h>
3
4 int shmget(key_t key, size_t size, int shmfig);
5
6 void *shmat(int shmid, void *shmaddr, int shmfig);
7
8 int shmdt(const void *shmaddr);
```

Shared memory and synchronization



- Reading and writing shared memory is typically a critical section
- Thus, we need to use sync. primitives to control access

1.1 Case Study: Data Transfer Between Processes

Case Study: Data Transfer Between Processes

- Two applications: a *writer* process and a *reader* process
- The writer process reads data from the keyboard (standard input) and places into the shared memory segment
- And the reader process reads the data from the shared memory segment and dumps it to the display (standard output)

```
1 #include <sys/types.h>
2 #include <sys/ipc.h>
3 #include <sys/shm.h>
4 #include <stdio.h>
5 #define SHMSZ 27
6 int main(int argc, char *argv[])
7 {
8     char c;
9     int shmid;
10    key_t key;
11    char *shm, *s;
12    // We'll name our shared memory segment
13    // "5678".
14    key = 5678;
15    // Create the segment.
```

```

17     if ((shmid = shmget(key, SHMSZ, IPC_CREAT | 0666)) < 0) {
18         exit(1);
19     }
20     // Now we attach the segment to our data space.
21     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
22         exit(1);
23     }
24     //
25     s = shm;
26     for (c = 'a'; c <= 'z'; c++)
27         *s++ = c;
28     *s = NULL;
29     //
30     // Finally, we wait until the other process
31     // changes the first character of our memory
32     // to '*', indicating that it has read what
33     // we put there.
34     while (*shm != '*')    sleep(1);
35     exit(0);

```

Shared Memory Example: Reader

```

1  #include <sys/types.h>
2  #include <sys/ipc.h>
3  #include <sys/shm.h>
4  #include <stdio.h>
5  #define SHMSZ    27
6  int main(int argc, char *argv[]) {
7      int shmid;
8      key_t key;
9      char *shm, *s;
10     // We need to get the segment named
11     // "5678", created by the server.
12     key = 5678;
13     // Locate the segment.
14     if ((shmid = shmget(key, SHMSZ, 0666)) < 0) {
15         exit(1);
16     }
17     // Now we attach the segment to our data space.
18     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
19         exit(1);
20     }
21     // Now read what the server put in the memory.
22     for (s = shm; *s != NULL; s++)
23         putchar(*s);
24     putchar('\n');
25     // Finally, change the first character of the
26     // segment to '*', indicating we have read
27     // the segment.
28     *shm = '*';
29     exit(0);
30 }

```

1.2 POSIX Shared Memory

System V Shared Memory Design Issues

- The SysV shared memory model uses keys and identifies to manage shared memory segments
 - This is inconsistent with the standard UNIX I/O model that uses filenames and file descriptors
- Use of shared file mappings for IPC requires the use of a disk file, even if the application doesn't require persistent backing store

The POSIX Shared Memory Model

- The POSIX model is a variation of “memory mapping” model
 - Pages of main memory are mapped into a virtual memory space
 - One may map a file (or other things) into this space and access it as if it was memory

The mmap() System Call

```
#include <sys/mman.h>
2
void *mmap(void *addr,
4         size_t length, int prot,
         int flags, int fd, off_t offset);
```

- The `addr` parameter indicates where in the virtual address space the mapped page should be placed.
 - If `NULL`, the OS chooses a suitable address in the memory map, this is preferred for portability issues

Example: Using memory mapping with files

```
1 #include <sys/mman.h>
  #include <fcntl.h>
3 #define MEMSIZE 10
  int main(int argc, char *argv[]) {
5     char *addr;
      int fd;
7     fd = open(argv[1], O_RDWR);
      if (fd == -1) exit(-1);
9     addr = mmap(NULL,
                  MEM_SIZE,
11                 PROT_READ | PROT_WRITE,
                  MAP_SHARED,
13                 fd,
                  0);
15     if (addr == MAP_FAILED) exit(-1);
      if (close(fd) == -1) /* No longer need 'fd' */
17         exit(-1);
      printf("Current string=.*s\n", MEM_SIZE, addr);
19     /* Secure practice: output at most MEM_SIZE bytes */
      memset(addr, 0, MEM_SIZE); /* Zero out region */
21     strncpy(addr, argv[2], MEM_SIZE - 1);
      if (msync(addr, MEM_SIZE, MS_SYNC) == -1)
23         exit(-1);
      printf("Copied \"%s\" to shared memory\n", argv[2]); }
25     exit(EXIT_SUCCESS);
}
```

Case Study: Message capture to log - Server

```
#include <stdio.h>
2 #include <stdlib.h>
#include <sys/types.h>
4 #include <sys/stat.h>
#include <fcntl.h>
6 #include <string.h>
#include <unistd.h>
8 #include <semaphore.h>
#include <sys/mman.h>
10 // Buffer data structures
#define MAX_BUFFERS 10
12 #define LOGFILE "/tmp/example.log"
#define SEM_MUTEX_NAME "/sem-mutex"
14 #define SEM_BUFFER_COUNT_NAME "/sem-buffer-count"
#define SEM_SPOOL_SIGNAL_NAME "/sem-spool-signal"
16 #define SHARED_MEM_NAME "/posix-shared-mem-example"

18 struct shared_memory {
    char buf [MAX_BUFFERS] [256];
20     int buffer_index;
    int buffer_print_index;
22 };

24 void error (char *msg);
// Print system error and exit
26 void error (char *msg)
{
28     perror (msg);
    exit (1);
30 }
```

Case Study: Message capture to log - Server

```
int main (int argc, char **argv)
2 {
    struct shared_memory *shared_mem_ptr;
4     sem_t *mutex_sem, *buffer_count_sem, *spool_signal_sem;
    int fd_shm, fd_log;
6     char mybuf [256];
    // Open log file
8     if ((fd_log = open (LOGFILE, O_CREAT | O_WRONLY | O_APPEND | O_SYNC, 0666)) == -1)
        error ("fopen");
10     // mutual exclusion semaphore, mutex_sem with an initial value 0.
    mutex_sem = sem_open (SEM_MUTEX_NAME, O_CREAT, 0660, 0);
12     if (mutex_sem == SEM_FAILED)
        error ("sem_open");
14     // Get shared memory
    fd_shm = shm_open (SHARED_MEM_NAME, O_RDWR | O_CREAT, 0660)
16     if (fd_shm == -1)
        error ("shm_open");
18     ftruncate (fd_shm, sizeof (struct shared_memory))
    if (fd_shm == -1)
20         error ("ftruncate");
```

Case Study: Message capture to log - Server

```

2 // Continuing from previous slide..
shared_mem_ptr = mmap (NULL, sizeof (struct shared_memory), PROT_READ | PROT_WRITE,
MAP_SHARED,
fd_shm, 0;
4 if (shared_mem_ptr == MAP_FAILED)
error ("mmap");
6 // Initialize the shared memory
shared_mem_ptr -> buffer_index = shared_mem_ptr -> buffer_print_index = NULL;
8
// counting semaphore, indicating the number of available
// buffers. Initial value = MAX_BUFFERS
10 buffer_count_sem = sem_open (SEM_BUFFER_COUNT_NAME, O_CREAT, 0660, MAX_BUFFERS);
12 if (buffer_count_sem == SEM_FAILED)
error ("sem_open");
14 // counting semaphore, indicating the number of strings to be
// printed. Initial value = 0
16 spool_signal_sem = sem_open (SEM_SPOOL_SIGNAL_NAME, O_CREAT, 0660, 0);
18 if (spool_signal_sem == SEM_FAILED)
error ("sem_open");

```

Case Study: Message capture to log - Server

```

2 // Initialization complete; now we can set mutex semaphore as 1 to
// indicate shared memory segment is available
if (sem_post (mutex_sem) == -1)
4 error ("sem_post: mutex_sem");
6 while (true) { // forever
// Is there a string to print? P (spool_signal_sem);
8 if (sem_wait (spool_signal_sem) == -1)
error ("sem_wait: spool_signal_sem");
10
strcpy(mybuf, shared_mem_ptr -> buf [shared_mem_ptr -> buffer_print_index]);
12
/* Since there is only one process (the logger) using the
buffer_print_index, mutex semaphore is not necessary */
14 (shared_mem_ptr -> buffer_print_index)++;
16 if (shared_mem_ptr -> buffer_print_index == MAX_BUFFERS)
shared_mem_ptr -> buffer_print_index = 0;
18
/* Contents of one buffer has been printed.
One more buffer is available for use by producers.
20 Release buffer: V (buffer_count_sem); */
22 if (sem_post(buffer_count_sem) == -1)
error ("sem_post: buffer_count_sem");
24
// write the string to file
26 if (write (fd_log, mybuf, strlen (mybuf)) != strlen (mybuf))
error ("write: logfile");
28 }
}

```

Case Study: Message capture to log - Client

```

1 int main (int argc, char **argv)
{
3 struct shared_memory *shared_mem_ptr;
sem_t *mutex_sem, *buffer_count_sem, *spool_signal_sem;

```

```

5  int fd_shm;
6  char mybuf [256];
7  struct shared_memory *shared_mem_ptr;
8  sem_t *mutex_sem, *buffer_count_sem, *spool_signal_sem;
9  int fd_shm, fd_log;
10 char mybuf [256];
11 // mutual exclusion semaphore, mutex_sem
12 mutex_sem = sem_open (SEM_MUTEX_NAME, 0, 0, 0);
13 if (mutex_sem == SEM_FAILED)
14     error ("sem_open");
15 // Get shared memory
16 fd_shm = shm_open (SHARED_MEM_NAME, O_RDWR, 0);
17 if (fd_shm == -1)
18     error ("shm_open");
19 shared_mem_ptr = mmap (NULL,
20                         sizeof (struct shared_memory),
21                         PROT_READ | PROT_WRITE, MAP_SHARED,
22                         fd_shm,
23                         0);
24 if (shared_mem_ptr == MAP_FAILED)
25     error ("mmap");

```

Case Study: Message capture to log - Client

```

1  // counting semaphore, indicating the number of available buffers.
2  buffer_count_sem = sem_open (SEM_BUFFER_COUNT_NAME, 0, 0, 0);
3  if (buffer_count_sem == SEM_FAILED)
4      error ("sem_open");
5  // counting semaphore, indicating the number of strings to be
6  // printed. Initial value = 0
7  spool_signal_sem = sem_open (SEM_SPOOL_SIGNAL_NAME, 0, 0, 0);
8  if (spool_signal_sem == SEM_FAILED)
9      error ("sem_open");

```

Case Study: Message capture to log - Client

```

1  char buf [200], *cp;
2  printf ("Please type a message: ");
3  while (fgets (buf, 198, stdin)) {
4      // remove newline from string
5      int length = strlen (buf);
6      if (buf [length - 1] == '\n')
7          buf [length - 1] = '\0';
8      // get a buffer: P (buffer_count_sem);
9      if (sem_wait (buffer_count_sem) == -1)
10         error ("sem_wait: buffer_count_sem");
11     /* There might be multiple producers. We must ensure that
12        only one producer uses buffer_index at a time. */
13     // P (mutex_sem);
14     if (sem_wait (mutex_sem) == -1)
15         error ("sem_wait: mutex_sem");

```

Case Study: Message capture to log - Client

```

1 // Critical section
2 time_t now = time (NULL);
3 cp = ctime (&now);
4 int len = strlen (cp);
5 if (*(cp + len - 1) == '\n')
6     *(cp + len - 1) = '\0';
7 sprintf (shared_mem_ptr -> buf [shared_mem_ptr -> buffer_index],
8         "%d: %s %s\n",
9         getpid (),
10        cp, buf);
11 (shared_mem_ptr -> buffer_index)++;
12 if (shared_mem_ptr -> buffer_index == MAX_BUFFERS)
13     shared_mem_ptr -> buffer_index = 0;
14 // Release mutex sem: V (mutex_sem)
15 if (sem_post (mutex_sem) == -1)
16     error ("sem_post: mutex_sem");
17 // Tell spooler that there is a string to print: V (spool_signal_sem);
18 if (sem_post (spool_signal_sem) == -1)
19     error ("sem_post: (spool_signal_sem)");
20 printf ("Please type a message: ");
21 }
22 if (munmap (shared_mem_ptr, sizeof (struct shared_memory)) == -1)
23     error ("munmap");
24 exit (0);
25 }

```

2 Pipes, Named and Otherwise

Pipes?

- Pipes are the oldest method of IPC on UNIX systems, having been defined in Classical UNIX in the early 1970s
- Solution to the question of “How can the shell allow the output produced by one process be used as input to a related process?”
- Named pipes (or FIFOs) are a variation on this idea that allows for communication between *any* processes

Pipes and the shell

- Command-line shells in UNIX-like OSes allow one to send the output from a command to the input of a second command:

```

1 ls | wc -l
find . -name "*.doc?" -print0 | xargs grep -i -n -e "test grade*"

```




Figure 44-1: Using a pipe to connect two processes

Characteristics of a pipe


- A pipe is a byte stream
- Reads from an empty pipe will block
- Reads from a closed write-end of a pipe will return EOF
- Pipes are unidirectional
- Writes of sizes up to the pipe capacity are guaranteed to be atomic
- Pipes have a limited capacity and writes will block until a reader pulls data from the pipe

pipe(): Opening a pipe

```
#include unistd.h>
2 int pipe(int filedescriptors[2]);
```

- A success call to `pipe()` returns two open file descriptors: `[0]` for the read-end of the pipe, `[1]` for the write end.
- Once opened, we treat the pipe as if it was any other type of file

Pipes, fork(), and exec()



- So, with multiple processes, one of the first things one does is reconfigure the ends of the pipe for the parent and child
- Parent is the writer, and so will close the read-end of the pipe
- Child the reader, and so will close the write-end of the pipe

Making things work in your code

```
int filedes[2];
2 if (pipe(filedes) == -1)
    error("pipe");
4 switch (fork()) {
5     case -1:
6         error("fork");
7     case 0: /* Child */
8         if (close(filedes[1]) == -1)
9             error("close");
10        /* Child now reads from pipe */
11        break;
12    default: /* Parent */
13        if (close(filedes[0]) == -1)
14            error("close");
15        /* Parent now writes to pipe */
16        break;
17 }
```

Remapping standard input or standard output

```
1 int pfd[2];  
  pipe(pfd);  
3 // Code from previous slide goes here to fork()  
  close(STDOUT_FILENO);  
5 dup(pfd[1]);
```

2.1 Named Pipes

A Named Pipe is a Pipe

- Also known as a FIFO (think about it)
- Main difference is that a named pipe has a name within the file system is opened in the same way as a regular file
- Named pipes are created using the **mkfifo** shell command
 - And a system call with the same name can be used to create a FIFO

Consider the following shell script:

```
1 mkfifo myfifo  
  wc -l < myfifo &  
3 ls -l | tee myfifo | sort -k5n
```